

The Knapsack Problem

Andrei RUSĂNESCU, Gabriel-Ioan PAVEL, Amir FALLAH-MIRZAEI, 323CC

University Politehnica of Bucharest

Abstract. The present paper analyses several algorithms that solve the Knapsack problem, taking into account both the speed and the memory used by the program. It presents three solutions, two of which always generate the correct answer, and one that obtains a result close to the right one.

1 Introduction

The Knapsack problem derives its name from the situation faced by someone who wants to fill their limited-capacity knapsack with the most valuable items that can fit. You can choose whether to put it in the bag or not put it at all, hence the name "0/1 Knapsack Problem". The problem has been studied for more than a century, with writings dating as far back as 1897.

2 Practical Applications

The Knapsack Problem and its variations have a multitude of real life applications including financial modeling, production and inventory management systems, stratified sampling, design of queuing network models in manufacturing, and control of traffic overload in telecommunication systems.

1. Financial Modeling

- **Investment Portfolio Selection:** In finance, the Knapsack Problem is used to select a combination of investment assets that maximize returns while staying within a budget or risk tolerance. This involves balancing the trade-offs between different financial instruments like stocks, bonds, and derivatives.
- **Capital Budgeting:** Companies use this problem to decide how to allocate limited capital resources among various potential projects to maximize returns.

2. Production and Inventory Management Systems

- **Optimal Resource Allocation:** In production, the problem helps in determining the best way to allocate limited resources like raw materials or machinery to different products to maximize output or profit.
- **Inventory Control:** It is applied to decide which items to stock in limited warehouse space, ensuring maximum profitability or utility while minimizing costs.

3. Stratified Sampling

- **Survey Design:** In statistics, the Knapsack Problem helps in stratified sampling by selecting a sample that represents different strata or subgroups of a population in a cost-effective way, ensuring that the sample is both diverse and within budget constraints.

4. Design of Queuing Network Models in Manufacturing

- **Production Line Optimization:** In manufacturing, the problem aids in the design and optimization of queuing networks, where the goal is to minimize waiting times and costs while maximizing throughput. This involves selecting the right combination of machines, processes and workflows within budgetary and space constraints.

5. Control of Traffic Overload in Telecommunication Systems

- **Bandwidth Allocation:** The Knapsack Problem is used to allocate limited bandwidth resources to various communication channels or services to maximize overall network efficiency and quality of service.
- **Data Packet Prioritization:** It helps in determining which data packets should be transmitted first based on their priority and the available network capacity to prevent congestion and optimize network performance.

3 NP-Hard Proof

The Knapsack problem has as inputs: **W** - the maximum weight of the sack, **N** - the number of objects and the **objects** themselves that have as proprieties: **weight and profit**. The program returns the maximum profit that you can get without exceeding the weight limit. In the following demonstration we only use the **fail()** NP function because it doesn't interrupt the execution of the other threads. We also assume that the assignments to the **MAX_PROFIT** variable are **atomic**.

```
MAX_PROFIT = 0 # the result will be kept here
def knapsack(W, N, Objects):
    seen = N * [0]
    current_weight = 0
    current_profit = 0
    sack = []

    # generation phase
    for i = 1 : N:
        c = choice(0...N)
        if c != 0:
            if seen[c] == 1:
                fail()
            seen[c] = 1
            sack.append(Objects[c])
            current_weight += Objects[c]['weight']
            if current_weight > W:
                fail()
            current_profit += Objects[c]['profit']

    # testing phase
    if current_profit > MAX_PROFIT:
        MAX_PROFIT = current_profit
        fail()

    fail()
```

4 Description of the algorithms

4.1 Dynamic programming

There are two approaches to solve this problem using dynamic programming: using a memoization table (top-down) or tabulation (bottom-up).

- **Memoization:** Using recursion to solve the Knapsack Problem results in an exponential time complexity. This is due to the fact that the algorithm recomputes functions that have already been processed multiple times. Instead of recomputing these values, we can utilize a *memoization* table to

store the results of these functions. By doing so, we can significantly reduce the time complexity and make the algorithm more efficient. Memoization involves storing the results of previously computed subproblems in a table (often a two-dimensional array).

- **Tabulation:** This is the approach I chose to implement in code, because it is usually faster than the memoization one owing to its iterative nature. Despite calculating all the combinations of objects in the knapsack, the tabulation method is known as "Bottom-Up" since it builds the solution incrementally from the smallest subproblems to the overall problem. Tabulation involves filling up a table (typically a two-dimensional array, even though it can be done using a one-dimensional array). Each entry in the table represents the solution to a subproblem, and the solution to the overall problem is derived from these subproblem solutions.

```
int N, G; // N - the number of objects, G - the capacity of the knapsack
std::vector<std::vector<int>> dp; // dp matrix
std::vector<int> weights, values; // weights and values of the objects
inline int knapsack() {
    for (int i = 0; i <= N; ++i)
        dp[i][0] = 0;
    for (int j = 0; j <= G; ++j)
        dp[0][j] = 0;

    for (int i = 1; i <= N; ++i) {
        for (int w = 1; w <= G; ++w) {
            if (weights[i - 1] <= w)
                dp[i][w] = std::max(values[i - 1] + dp[i - 1][w - weights[i - 1]], dp[i - 1][w]);
            else
                dp[i][w] = dp[i - 1][w];
        }
    }
    return dp[N][G];
}

inline std::vector<int> find_objects() {
    std::vector<int> res;
    int m = N, n = G;
    while (n > 0 && m > 0) {
        if (dp[m][n] != dp[m - 1][n]) {
            n -= weights[m - 1];
            res.push_back(m);
        }
        --m;
    }
    std::reverse(res.begin(), res.end());
    return res;
}
```

To begin with, the code is as fast as it can possibly be, while still being intelligible. I used inlining to suggest the compiler to paste this code where it is called from in order to reduce the potential overhead of pushing registers onto the stack. Furthermore, I used global variables that are initialized in the main procedure, due to the fact that these two functions are benchmarked using the Google Benchmark for C++. I wanted to make a separation between I/O (filling the data), allocations and initializations, and the true computing effort these functions bring.

Each row in the matrix corresponds to an object, while each column represents a maximum weight capacity. The entry $dp[i][w]$ denotes the maximum profit achievable using the first i objects with a weight limit of w . This value is determined by taking the maximum between two options: the sum of the current object's profit and the optimal profit from the remaining capacity, and the best profit obtained without including the current object. In essence, it captures the overall maximum profit for the given weight constraint, the last cell being the result. The function responsible for finding the objects that were put in the bag is doing reverse engineering on the matrix. If the profit obtained for the current row is different than the one above it, meaning the current object was used, so it is selected, otherwise, the current value is originating from an upper row.

Complexity Analysis

Time Complexity The algorithm involves nested loops to fill the dynamic programming (DP) table: The outer loop runs from 1 to N (the number of objects), resulting in N iterations. The inner loop runs from 1 to G (the capacity of the knapsack), resulting in G iterations. For each combination of i (object) and w (weight), a constant amount of work is done, either updating the DP table or performing a comparison. Thus, the overall time complexity of the algorithm is $\mathcal{O}(N * G)$.

Space Complexity The space complexity is determined by the DP table, which is a two-dimensional vector of size $(N+1)*(G+1)$. This results in a space complexity of $\mathcal{O}(N * G)$. Additionally, the *find objects* function uses a vector to store the selected objects, which in the worst case, may require space proportional to N , but this does not change the overall space complexity.

Advantages and Disadvantages

Advantages The main advantages are the speed of execution, as the result can be obtained in a reasonable amount of time, and the avoidance of issues that might arise with recursion, such as stack overflow.

Disadvantages The main disadvantages include the large memory usage due to the DP table and the fact that the matrix size cannot be dynamically adjusted without recalculating the values.

4.2 Backtracking

This approach explores every possible combination of objects recursively. The algorithm is initiated by the main procedure, which sets up the necessary variables and calls the main backtracking function. In this context, **G** represents the maximum weight the knapsack can carry, **maxx** is the maximum profit found so far. The process concludes by returning the best possible result. The selection of the indexes is done when a subset of N elements has been explored and its profit is the current maximal value.

```
int N, G;
int maxx;
std::vector<int> weights, values;
std::vector<int> selection, current_selection;

int knapsack(int index, int curr_weight, int curr_profit) {
    if (index == N) {
        if (curr_profit > maxx) {
            maxx = curr_profit;
            selection = current_selection;
        }
        return maxx;
    }

    if (curr_weight + weights[index] <= G) {
        current_selection.push_back(index + 1);
        maxx = std::max(maxx, knapsack(index + 1, curr_weight +
            weights[index], curr_profit + values[index]));
        current_selection.pop_back();
    }
    return std::max(maxx, knapsack(index + 1, curr_weight,
        curr_profit));
}
```

For each function call, the algorithm checks if the current object can be added without exceeding the maximum weight. If it can, the function updates the current maximum profit and continues to explore deeper. If not, it skips adding the object. The recursion stops when all combinations of objects have been considered.

The function iterates over all objects starting from a given **index**. For each object, it adds its weight and value to the current totals, then dives deeper into the recursion. After exploring that path, it removes the object's weight and value to backtrack and explore other possibilities.

Complexity Analysis

Time Complexity The main factor contributing to the time complexity is the exhaustive nature of the algorithm, it explores every possible subset of objects.

This results in 2^n iterations, as each object can either be included or excluded. Therefore, the time complexity is $\mathcal{O}(2^n)$, which grows exponentially with the number of objects.

Space Complexity The space complexity is driven by the depth of the recursion stack. With n objects to check, the stack's depth can reach up to n , making the space complexity $\mathcal{O}(n)$.

Advantages and Disadvantages

Advantages One of the main advantages of the backtracking approach is its simplicity. The algorithm is straightforward to implement and understand, making it an excellent choice for educational purposes or scenarios where optimality is required for small inputs. Additionally, it doesn't require significant memory beyond the recursion stack and a few auxiliary variables.

Disadvantages The primary drawback is the method's inefficiency due to its brute-force nature. The exponential time complexity makes it impractical for large input sizes, as it results in a vast number of redundant calculations. This inefficiency becomes a major limitation, especially when compared to more optimized approaches like dynamic programming.

4.3 Greedy

To solve the problem using an algorithm that provides a result close to the correct one I used a greedy approach. It solves the fractional knapsack problem. The algorithm uses a heuristic based on the value-to-weight ratio of items to make decisions. Basically, it sorts the objects based on the value to weight ratio and then fills the bag with the objects until the last one that is divided. If the sorting of the objects somehow is the correct order, the objects are not truncated at all. The overall value is close to the correct one based on a large number of tests.

```
int N, G;
std::vector<std::vector<int>>> pairs;
std::vector<int> selection;

struct fcmp {
    bool operator()(const std::vector<int> &a, const std::vector<int>
        &b) {
        double r1 = (double)a[1] / (double)a[0];
        double r2 = (double)b[1] / (double)b[0];
        return r1 > r2;
    }
};

double knapsack() {
```

```

std::sort(pairs.begin(), pairs.end(), fcmp());
int current_weight = G;
double res = 0;
for (int i = 0; i < N; ++i) {
    if (pairs[i][0] <= current_weight) {
        current_weight -= pairs[i][0];
        selection.push_back(pairs[i][2]);
        res += pairs[i][1];
    } else {
        selection.push_back(pairs[i][2]);
        res += pairs[i][1] * ((double)current_weight) /
            (double)pairs[i][0];
        break;
    }
}

return res;
}

```

I used a vector of objects that have 3 values, the weight, the value and the corresponding index. For sorting, I used struct `fcmp` to provide a sorting logic for the `sort()` function.

Complexity Analysis

Time Complexity The time complexity of the algorithm is $\mathcal{O}(N \log N)$ due to the sorting step, where N is the number of items. After sorting, the algorithm iterates through the list of items, which takes $\mathcal{O}(N)$ time. Thus, the overall time complexity is dominated by the sorting step, resulting in $\mathcal{O}(N \log N)$.

Space Complexity The space complexity is $\mathcal{O}(N)$ because the algorithm uses additional space to store the list of items and the selection vector.

Advantages and Disadvantages

Advantages The algorithm is efficient and easy to implement, with a relatively low time complexity of $\mathcal{O}(N \log N)$. It provides a quick approximation for the knapsack problem, especially useful when an exact solution is computationally expensive. The greedy approach ensures that the solution is close to optimal for the fractional knapsack problem.

Disadvantages The algorithm may not provide the exact optimal solution for the 0/1 knapsack problem, where items cannot be divided. It relies on a heuristic (*value-to-weight* ratio) that may not always yield the best solution for specific instances of the problem. The solution's accuracy depends on the distribution of item weights and values; in cases where high-value items have low weight, the algorithm performs well, but it may underperform in other cases.

5 Testing and Evaluation

The tests are 30 in number and were partly generated using a python script that generated tests with sizes between 20 and 100 and weights randomly chosen between 500 and 1000. The other part is made of 4 tests I chose from an online resource I cited in the bibliography. To run the benchmark, you should execute the `run_benchmarks.sh` script and it will write the results in the tests/benchmarks folder.

For the benchmark test I used my personal laptop with an AMD Ryzen 7 3750H with Radeon Vega Mobile Gfx 2.30 GHz, 16.0 GB of RAM (13.9 GB usable), running Ubuntu 22.04. The tests were run using the gcc compiler, version gcc version 11.4.0 with the flags: `-std=c++11 -isystem utils/benchmark/include -L../utils/benchmark/build/src -lbenchmark -lpthread`. All the tests were run with the laptop plugged in, on the Turbo mode from Ubuntu.

The results are summarized in the tables below, presenting the performance of each algorithm. The number of *iterations* for each benchmark test is determined automatically by the benchmarking tool, which ensures that the functions are called a sufficient number of times to provide reliable measurements. Each line corresponds to a test. The tests are located in the tests/input folder.

5.1 Performance Analysis

1. **Greedy Algorithm:** Speed: The Greedy algorithm consistently outperforms the other two in terms of speed. This efficiency is due to its simplistic approach, which just sorts the items and then fills the bag. Complexity: The Greedy algorithm operates with a time complexity of $\mathcal{O}(N \log N)$. However, it does not guarantee the optimal solution for the problem. Use Case: Best suited for scenarios where speed is crucial, and the problem constraints or inputs naturally lend themselves to a near-optimal solution, close to the correct one.

2. **Dynamic Programming (DP):** Speed: The DP approach shows moderate performance, sitting between the Greedy and Backtracking methods. This is expected given its polynomial time complexity. Complexity: With a time complexity of $\mathcal{O}(N * W)$, where N is the number of items and W is the maximum weight capacity, DP provides a balance between optimality and efficiency. It ensures the optimal solution by considering all possible combinations of items up to the maximum weight. Use Case: Ideal for scenarios where an exact solution is necessary.

3. **Backtracking Algorithm:** Speed: Backtracking is significantly slower, particularly as the input size increases. This is evident from the exponential increase in execution time as the value of N grows. Complexity: The time complexity is $\mathcal{O}(2^N)$, making it impractical for large tests. Use Case: While not efficient for large inputs, backtracking can be useful in cases where exact solutions are needed for small tests.

Benchmark DP	Time (ns)	CPU (ns)	Iterations
knapsack_handler_tabulation	397274	397277	1851
knapsack_handler_tabulation	451001	450761	1728
knapsack_handler_tabulation	381943	381887	1507
knapsack_handler_tabulation	489491	489492	1326
knapsack_handler_tabulation	493156	493152	1298
knapsack_handler_tabulation	435406	435404	1636
knapsack_handler_tabulation	650042	650031	917
knapsack_handler_tabulation	532384	532373	1252
knapsack_handler_tabulation	516264	516259	1470
knapsack_handler_tabulation	562368	562380	1237
knapsack_handler_tabulation	883800	883352	735
knapsack_handler_tabulation	473825	473815	1522
knapsack_handler_tabulation	701742	701745	875
knapsack_handler_tabulation	639519	639524	1058
knapsack_handler_tabulation	962106	962057	658
knapsack_handler_tabulation	929154	929151	827
knapsack_handler_tabulation	698738	698732	1040
knapsack_handler_tabulation	793952	793950	869
knapsack_handler_tabulation	902861	902869	714
knapsack_handler_tabulation	624316	624326	1126
knapsack_handler_tabulation	1089568	1089519	594
knapsack_handler_tabulation	843631	843624	876
knapsack_handler_tabulation	884242	884247	717
knapsack_handler_tabulation	729051	728534	866
knapsack_handler_tabulation	1303246	1303232	550
knapsack_handler_tabulation	1288841	1288492	539
knapsack_handler_tabulation	741127	741112	834
knapsack_handler_tabulation	1092041	1092014	606
knapsack_handler_tabulation	399632	399628	1952
knapsack_handler_tabulation	294332	294297	2375

Table 1. Benchmark Results for knapsack_handler_tabulation

Benchmark	Backtracking	Time (ns)	CPU (ns)	Iterations
knapsack_handler_backtracking		422676	422635	1644
knapsack_handler_backtracking		1429100	1429023	483
knapsack_handler_backtracking		6927283	6925875	103
knapsack_handler_backtracking		11765128	11765176	58
knapsack_handler_backtracking		6308162	6308182	111
knapsack_handler_backtracking		2075346	2075331	334
knapsack_handler_backtracking		23259420	23259193	30
knapsack_handler_backtracking		2290172	2290153	311
knapsack_handler_backtracking		31302150	31301750	24
knapsack_handler_backtracking		163084480	163085075	4
knapsack_handler_backtracking		559927626	559929100	1
knapsack_handler_backtracking		1736294	1736248	410
knapsack_handler_backtracking		18161163	18160888	40
knapsack_handler_backtracking		89581115	89579544	9
knapsack_handler_backtracking		3263062506	3263021100	1
knapsack_handler_backtracking		3040190737	3040169900	1
knapsack_handler_backtracking		211758763	211757767	3
knapsack_handler_backtracking		624621926	624091200	1
knapsack_handler_backtracking		1359326241	1359322300	1
knapsack_handler_backtracking		72809409	72808410	10
knapsack_handler_backtracking		6523479592	6523417800	1
knapsack_handler_backtracking		342279551	342120750	2
knapsack_handler_backtracking		1379948451	1379898500	1
knapsack_handler_backtracking		125158735	125157260	5
knapsack_handler_backtracking		5962969626	5962870200	1
knapsack_handler_backtracking		34420000000	34420000000	1
knapsack_handler_backtracking		123940051	123939083	6
knapsack_handler_backtracking		31164000000	31163000000	1
knapsack_handler_backtracking		47379010	47379250	14
knapsack_handler_backtracking		2652311691	2652282900	1

Table 2. Benchmark Results for knapsack_handler_backtracking

Benchmark Greedy	Time (ns)	CPU (ns)	Iterations
knapsack_handler_greedy	7195	7193	93873
knapsack_handler_greedy	8058	8058	86586
knapsack_handler_greedy	9510	9510	77823
knapsack_handler_greedy	9905	9905	64998
knapsack_handler_greedy	9600	9600	72680
knapsack_handler_greedy	7909	7909	89973
knapsack_handler_greedy	13484	13484	52825
knapsack_handler_greedy	15715	15715	48553
knapsack_handler_greedy	16990	16990	45755
knapsack_handler_greedy	16546	16545	42428
knapsack_handler_greedy	18386	18366	34185
knapsack_handler_greedy	17087	17087	39253
knapsack_handler_greedy	18532	18528	38958
knapsack_handler_greedy	18495	18495	38447
knapsack_handler_greedy	20672	20671	32829
knapsack_handler_greedy	21150	21148	33215
knapsack_handler_greedy	21175	21175	33052
knapsack_handler_greedy	21954	21954	28701
knapsack_handler_greedy	23307	23307	34267
knapsack_handler_greedy	19612	19612	29183
knapsack_handler_greedy	21941	21938	36639
knapsack_handler_greedy	26130	26130	26768
knapsack_handler_greedy	28771	28741	24319
knapsack_handler_greedy	30160	30160	21867
knapsack_handler_greedy	30992	30991	22807
knapsack_handler_greedy	31482	31482	21237
knapsack_handler_greedy	34720	34719	21059
knapsack_handler_greedy	34224	34224	20451
knapsack_handler_greedy	7058	7058	90852
knapsack_handler_greedy	10121	10121	71612

Table 3. Benchmark Results for knapsack_handler_greedy

6 Conclusions

The following images are graphs representing the performance of each algorithm:




Fig. 1. Time taken for the Greedy benchmarks




Fig. 2. Time taken for the Backtracking benchmarks



Fig. 3. Time taken for the Tabulation benchmarks

The graphs show the distribution of the performances of the algorithms based on the number of inputs and the execution time. Every algorithm shows an ascending trend, meaning that, generally, they take longer to complete the higher the number of sets of inputs they are given.

Some tests appear to be faster than others that have less sets of inputs. This is predominantly happening because the benchmarks were run in a multithreaded environment, where context switches were happening.

It is easy to say that the tabulation approach is by far the most reliable, albeit not the fastest. The speed of the greedy approach can be deceiving, suggesting it would be the best algorithm to solve the problem, but we have to take into consideration the fact that it does not provide the highest possible profit. Its solution is *almost* correct. With all of that in mind, if we are not looking for the perfect combination of items, but for a quick and decent result, the greedy would be the best algorithm. As seen in the second graph, the backtracking approach is by far the slowest of them all, since it explores every possible combination of items (time complexity of $\mathcal{O}(2^n)$). Even though it is very slow, it is still better in terms of the correctness of the result than the greedy algorithm. The only point it actually excels in is its simplicity of the implementation, as it's a good starting point to understand the problem.

With all of that being said, if we care about achieving a perfect result in a timely manner, we can conclude that the dynamic programming approach is the best one, and if we want a solution as fast as possible, although not perfect, the greedy algorithm is the way to go.

References

1. *Knapsack problem*, https://en.wikipedia.org/wiki/Knapsack_problem Accessed at [04/01/2025].
2. Kayode Badiru: *KNAPSACK PROBLEMS; METHODS, MODELS AND APPLICATIONS*, <https://shareok.org/server/api/core/bitstreams/7bb97e22-4c5d-44b3-b756-74aba3fc1fd7/content> Accessed at [05/01/2025]
3. LNCS Homepage, <http://www.springer.com/lncs>, Accessed at [04/01/2025]
4. *A Phase Angle-Modulated Bat Algorithm with Application to Antenna Topology Optimization* https://www.researchgate.net/publication/349878676_A_Phase_Angle-Modulated_Bat_Algorithm_with_Application_to_Antenna_Topology_Optimization, Accessed at [04/01/2025]
5. Geeks for Geeks, *0/1 Knapsack Problem*, <https://www.geeksforgeeks.org/0-1-knapsack-problem-dp-10/>, Accessed at [05/01/2025]