

The Knapsack Problem

Andrei RUSĂNESCU, Gabriel-Ioan PAVEL, Amir FALLAH-MIRZAEI, 323CC

University Politehnica of Bucharest

Abstract. The present paper analyses several algorithms that solve the Knapsack problem, taking into account both the speed and the memory used by the program. It presents three solutions, two of which always generate the correct answer, and one that obtains a result close to the right one.

1 Introduction

The Knapsack problem derives its name from the situation faced by someone who wants to fill their limited-capacity knapsack with the most valuable items that can fit. You can choose whether to put it in the bag or not put it at all, hence the name "0/1 Knapsack Problem". The problem has been studied for more than a century, with writings dating as far back as 1897.

2 Practical Applications

The Knapsack Problem and its variations have a multitude of real life applications including financial modeling, production and inventory management systems, stratified sampling, design of queuing network models in manufacturing, and control of traffic overload in telecommunication systems.

1. Financial Modeling

- **Investment Portfolio Selection:** In finance, the Knapsack Problem is used to select a combination of investment assets that maximize returns while staying within a budget or risk tolerance. This involves balancing the trade-offs between different financial instruments like stocks, bonds, and derivatives.
- **Capital Budgeting:** Companies use this problem to decide how to allocate limited capital resources among various potential projects to maximize returns.

2. Production and Inventory Management Systems

- **Optimal Resource Allocation:** In production, the problem helps in determining the best way to allocate limited resources like raw materials or machinery to different products to maximize output or profit.
- **Inventory Control:** It is applied to decide which items to stock in limited warehouse space, ensuring maximum profitability or utility while minimizing costs.

3. Stratified Sampling

- **Survey Design:** In statistics, the Knapsack Problem helps in stratified sampling by selecting a sample that represents different strata or subgroups of a population in a cost-effective way, ensuring that the sample is both diverse and within budget constraints.

4. Design of Queuing Network Models in Manufacturing

- **Production Line Optimization:** In manufacturing, the problem aids in the design and optimization of queuing networks, where the goal is to minimize waiting times and costs while maximizing throughput. This involves selecting the right combination of machines, processes and workflows within budgetary and space constraints.

5. Control of Traffic Overload in Telecommunication Systems

- **Bandwidth Allocation:** The Knapsack Problem is used to allocate limited bandwidth resources to various communication channels or services to maximize overall network efficiency and quality of service.
- **Data Packet Prioritization:** It helps in determining which data packets should be transmitted first based on their priority and the available network capacity to prevent congestion and optimize network performance.

3 NP-Hard Proof

4 Description of the algorithms

4.1 Dynamic programming

There are two approaches to solve this problem using dynamic programming: using a memoization table (top-down) or tabulation (bottom-up).

- **Memoization:** Using recursion to solve the Knapsack Problem results in an exponential time complexity. This is due to the fact that the algorithm recomputes functions that have already been processed multiple times. Instead of recomputing these values, we can utilize a *memoization* table to store the results of these functions. By doing so, we can significantly reduce the time complexity and make the algorithm more efficient. Memoization involves storing the results of previously computed subproblems in a table (often a two-dimensional array).
- **Tabulation:** This is the approach I chose to implement in code, because it is usually faster than the memoization one owing to its iterative nature. Despite calculating all the combinations of objects in the knapsack, the tabulation method is known as "Bottom-Up" since it builds the solution incrementally from the smallest subproblems to the overall problem. Tabulation involves

filling up a table (typically a two-dimensional array, even though it can be done using a one-dimensional array). Each entry in the table represents the solution to a subproblem, and the solution to the overall problem is derived from these subproblem solutions.

```
int N, G; // N - the number of objects, G - the capacity of the knapsack
std::vector<std::vector<int>> dp; // dp matrix
std::vector<int> weights, values; // weights and values of the objects
inline int knapsack() {
    for (int i = 0; i <= N; ++i)
        dp[i][0] = 0;
    for (int j = 0; j <= G; ++j)
        dp[0][j] = 0;

    for (int i = 1; i <= N; ++i) {
        for (int w = 1; w <= G; ++w) {
            if (weights[i - 1] <= w)
                dp[i][w] = std::max(values[i - 1] + dp[i - 1][w - weights[i - 1]], dp[i - 1][w]);
            else
                dp[i][w] = dp[i - 1][w];
        }
    }
    return dp[N][G];
}

inline std::vector<int> find_objects() {
    std::vector<int> res;
    int m = N, n = G;
    while (n > 0 && m > 0) {
        if (dp[m][n] != dp[m - 1][n]) {
            n -= weights[m - 1];
            res.push_back(m);
        }
        --m;
    }
    std::reverse(res.begin(), res.end());
    return res;
}
```

To begin with, the code is as fast as it can possibly be, while still being intelligible. I used inlining to suggest the compiler to paste this code where it is called from in order to reduce the potential overhead of pushing registers onto the stack. Furthermore, I used global variables that are initialized in the main procedure, due to the fact that these two functions are benchmarked using the Google Benchmark for C++. I wanted to make a separation between I/O (filling the data), allocations and initializations, and the true computing effort these functions bring.

Each row in the matrix corresponds to an object, while each column represents a maximum weight capacity. The entry $dp[i][w]$ denotes the maximum profit achievable using the first i objects with a weight limit of w . This value is determined by taking the maximum between two options: the sum of the current object's profit and the optimal profit from the remaining capacity, and the best profit obtained without including the current object. In essence, it captures the overall maximum profit for the given weight constraint, the last cell being the result. The function responsible for finding the objects that were put in the bag is doing reverse engineering on the matrix. If the profit obtained for the current row is different than the one above it, meaning the current object was used, so it is selected, otherwise, the current value is originating from an upper row.

Complexity Analysis

Time Complexity The algorithm involves nested loops to fill the dynamic programming (DP) table: The outer loop runs from 1 to N (the number of objects), resulting in N iterations. The inner loop runs from 1 to G (the capacity of the knapsack), resulting in G iterations. For each combination of i (object) and w (weight), a constant amount of work is done, either updating the DP table or performing a comparison. Thus, the overall time complexity of the algorithm is $O(N * G)$.

Space Complexity The space complexity is determined by the DP table, which is a two-dimensional vector of size $(N+1)*(G+1)$. This results in a space complexity of $O(N * G)$. Additionally, the *find objects* function uses a vector to store the selected objects, which in the worst case, may require space proportional to N , but this does not change the overall space complexity.

Advantages and Disadvantages

Advantages The main advantages are the speed of execution, as the result can be obtained in a reasonable amount of time, and the avoidance of issues that might arise with recursion, such as stack overflow.

Disadvantages The main disadvantages include the large memory usage due to the DP table and the fact that the matrix size cannot be dynamically adjusted without recalculating the values.

4.2 Backtracking

This approach explores every possible combination of objects recursively. The algorithm is initiated by the main procedure, which sets up the necessary variables and calls the main backtracking function. In this context, G represents the maximum weight the knapsack can carry, \mathbf{maxx} is the maximum profit found so far. The process concludes by returning the best possible result. The selection

of the indexes is done when a subset of N elements has been explored and its profit is the current maximal value.

```
int N, G;
int maxx;
std::vector<int> weights, values;
std::vector<int> selection, current_selection;

int knapsack(int index, int curr_weight, int curr_profit) {
    if (index == N) {
        if (curr_profit > maxx) {
            maxx = curr_profit;
            selection = current_selection;
        }
        return maxx;
    }

    if (curr_weight + weights[index] <= G) {
        current_selection.push_back(index + 1);
        maxx = std::max(maxx, knapsack(index + 1, curr_weight +
            weights[index], curr_profit + values[index]));
        current_selection.pop_back();
    }
    return std::max(maxx, knapsack(index + 1, curr_weight, curr_profit));
}
```

For each function call, the algorithm checks if the current object can be added without exceeding the maximum weight. If it can, the function updates the current maximum profit and continues to explore deeper. If not, it skips adding the object. The recursion stops when all combinations of objects have been considered.

The function iterates over all objects starting from a given **index**. For each object, it adds its weight and value to the current totals, then dives deeper into the recursion. After exploring that path, it removes the object's weight and value to backtrack and explore other possibilities.

Complexity Analysis

Time Complexity The main factor contributing to the time complexity is the exhaustive nature of the algorithm, it explores every possible subset of objects. This results in 2^n iterations, as each object can either be included or excluded. Therefore, the time complexity is $\mathcal{O}(2^n)$, which grows exponentially with the number of objects.

Space Complexity The space complexity is driven by the depth of the recursion stack. With n objects to check, the stack's depth can reach up to n , making the space complexity $\mathcal{O}(n)$.

Advantages and Disadvantages

Advantages One of the main advantages of the backtracking approach is its simplicity. The algorithm is straightforward to implement and understand, making it an excellent choice for educational purposes or scenarios where optimality is required for small inputs. Additionally, it doesn't require significant memory beyond the recursion stack and a few auxiliary variables.

Disadvantages The primary drawback is the method's inefficiency due to its brute-force nature. The exponential time complexity makes it impractical for large input sizes, as it results in a vast number of redundant calculations. This inefficiency becomes a major limitation, especially when compared to more optimized approaches like dynamic programming.

4.3 Greedy

To solve the problem using an algorithm that provides a result close to the correct one I used a greedy approach. It solves the fractional knapsack problem. The algorithm uses a heuristic based on the value-to-weight ratio of items to make decisions. Basically, it sorts the objects based on the value to weight ratio and then fills the bag with the objects until the last one that is divided. If the sorting of the objects somehow is the correct order, the objects are not truncated at all. The overall value is close to the correct one based on a large number of tests.

```
int N, G;
std::vector<std::vector<int>> pairs;
std::vector<int> selection;

struct fcmp {
    bool operator()(const std::vector<int> &a, const std::vector<int>
        &b) {
        double r1 = (double)a[1] / (double)a[0];
        double r2 = (double)b[1] / (double)b[0];
        return r1 > r2;
    }
};

double knapsack() {
    std::sort(pairs.begin(), pairs.end(), fcmp());
    int current_weight = G;
    double res = 0;
    for (int i = 0; i < N; ++i) {
        if (pairs[i][0] <= current_weight) {
            current_weight -= pairs[i][0];
            selection.push_back(pairs[i][2]);
            res += pairs[i][1];
        } else {
```

```

        selection.push_back(pairs[i][2]);
        res += pairs[i][1] * ((double)current_weight) /
            (double)pairs[i][0];
        break;
    }
}

return res;
}

```

I used a vector of objects that have 3 values, the weight, the value and the corresponding index. For sorting, I used struct fcmp to provide a sorting logic for the sort() function.

Complexity Analysis

Time Complexity The time complexity of the algorithm is $\mathcal{O}(N \log N)$ due to the sorting step, where N is the number of items. After sorting, the algorithm iterates through the list of items, which takes $\mathcal{O}(N)$ time. Thus, the overall time complexity is dominated by the sorting step, resulting in $\mathcal{O}(N \log N)$.

Space Complexity The space complexity is $\mathcal{O}(N)$ because the algorithm uses additional space to store the list of items and the selection vector.

Advantages and Disadvantages

Advantages The algorithm is efficient and easy to implement, with a relatively low time complexity of $\mathcal{O}(N \log N)$. It provides a quick approximation for the knapsack problem, especially useful when an exact solution is computationally expensive. The greedy approach ensures that the solution is close to optimal for the fractional knapsack problem.

Disadvantages The algorithm may not provide the exact optimal solution for the 0/1 knapsack problem, where items cannot be divided. It relies on a heuristic (*value-to-weight* ratio) that may not always yield the best solution for specific instances of the problem. The solution's accuracy depends on the distribution of item weights and values; in cases where high-value items have low weight, the algorithm performs well, but it may underperform in other cases.

5 Testing and Evaluation

The tests are 30 in number and were partly generated using a python script that generated tests with sizes between 20 and 100 and weights randomly chosen between 500 and 1000. The other part is made of 4 tests I chose from an online resource I cited in the bibliography. To run the benchmark, you should execute

the `run_benchmarks.sh` script and it will write the results in the `tests/benchmarks` folder.

For the benchmark test I used my personal laptop with an AMD Ryzen 7 3750H with Radeon Vega Mobile Gfx 2.30 GHz, 16.0 GB of RAM (13.9 GB usable), running Ubuntu 22.04. The tests were run using the gcc compiler, version gcc version 11.4.0 with the flags: `-std=c++11 -isystem utils/benchmark/include -L../utils/benchmark/build/src -lbenchmark -lpthread`. All the tests were run with the laptop plugged in, on the Turbo mode from Ubuntu.

The results are summarized in the tables below, presenting the performance of each algorithm. The number of *iterations* for each benchmark test is determined automatically by the benchmarking tool, which ensures that the functions are called a sufficient number of times to provide reliable measurements. Each line corresponds to a test. The tests are located in the `tests/input` folder.

5.1 Performance Analysis

1. **Greedy** Algorithm: Speed: The Greedy algorithm consistently outperforms the other two in terms of speed. This efficiency is due to its simplistic approach, which just sorts the items and then fills the bag. Complexity: The Greedy algorithm operates with a time complexity of $\mathcal{O}(N \log N)$. However, it does not guarantee the optimal solution for the problem. Use Case: Best suited for scenarios where speed is crucial, and the problem constraints or inputs naturally lend themselves to a near-optimal solution, close to the correct one.

2. **Dynamic Programming (DP)**: Speed: The DP approach shows moderate performance, sitting between the Greedy and Backtracking methods. This is expected given its polynomial time complexity. Complexity: With a time complexity of $\mathcal{O}(N * W)$, where N is the number of items and W is the maximum weight capacity, DP provides a balance between optimality and efficiency. It ensures the optimal solution by considering all possible combinations of items up to the maximum weight. Use Case: Ideal for scenarios where an exact solution is necessary.

3. **Backtracking** Algorithm: Speed: Backtracking is significantly slower, particularly as the input size increases. This is evident from the exponential increase in execution time as the value of N grows. Complexity: The time complexity is $\mathcal{O}(2^N)$, making it impractical for large tests. Use Case: While not efficient for large inputs, backtracking can be useful in cases where exact solutions are needed for small tests.

| Benchmark DP | Time (ns) | CPU (ns) | Iterations |
|-----------------------------|-----------|----------|------------|
| knapsack_handler_tabulation | 397274 | 397277 | 1851 |
| knapsack_handler_tabulation | 451001 | 450761 | 1728 |
| knapsack_handler_tabulation | 381943 | 381887 | 1507 |
| knapsack_handler_tabulation | 489491 | 489492 | 1326 |
| knapsack_handler_tabulation | 493156 | 493152 | 1298 |
| knapsack_handler_tabulation | 435406 | 435404 | 1636 |
| knapsack_handler_tabulation | 650042 | 650031 | 917 |
| knapsack_handler_tabulation | 532384 | 532373 | 1252 |
| knapsack_handler_tabulation | 516264 | 516259 | 1470 |
| knapsack_handler_tabulation | 562368 | 562380 | 1237 |
| knapsack_handler_tabulation | 883800 | 883352 | 735 |
| knapsack_handler_tabulation | 473825 | 473815 | 1522 |
| knapsack_handler_tabulation | 701742 | 701745 | 875 |
| knapsack_handler_tabulation | 639519 | 639524 | 1058 |
| knapsack_handler_tabulation | 962106 | 962057 | 658 |
| knapsack_handler_tabulation | 929154 | 929151 | 827 |
| knapsack_handler_tabulation | 698738 | 698732 | 1040 |
| knapsack_handler_tabulation | 793952 | 793950 | 869 |
| knapsack_handler_tabulation | 902861 | 902869 | 714 |
| knapsack_handler_tabulation | 624316 | 624326 | 1126 |
| knapsack_handler_tabulation | 1089568 | 1089519 | 594 |
| knapsack_handler_tabulation | 843631 | 843624 | 876 |
| knapsack_handler_tabulation | 884242 | 884247 | 717 |
| knapsack_handler_tabulation | 729051 | 728534 | 866 |
| knapsack_handler_tabulation | 1303246 | 1303232 | 550 |
| knapsack_handler_tabulation | 1288841 | 1288492 | 539 |
| knapsack_handler_tabulation | 741127 | 741112 | 834 |
| knapsack_handler_tabulation | 1092041 | 1092014 | 606 |
| knapsack_handler_tabulation | 399632 | 399628 | 1952 |
| knapsack_handler_tabulation | 294332 | 294297 | 2375 |

Table 1. Benchmark Results for knapsack_handler_tabulation

| Benchmark | Backtracking | Time (ns) | CPU (ns) | Iterations |
|-------------------------------|---------------------|------------------|-----------------|-------------------|
| knapsack_handler_backtracking | | 422676 | 422635 | 1644 |
| knapsack_handler_backtracking | | 1429100 | 1429023 | 483 |
| knapsack_handler_backtracking | | 6927283 | 6925875 | 103 |
| knapsack_handler_backtracking | | 11765128 | 11765176 | 58 |
| knapsack_handler_backtracking | | 6308162 | 6308182 | 111 |
| knapsack_handler_backtracking | | 2075346 | 2075331 | 334 |
| knapsack_handler_backtracking | | 23259420 | 23259193 | 30 |
| knapsack_handler_backtracking | | 2290172 | 2290153 | 311 |
| knapsack_handler_backtracking | | 31302150 | 31301750 | 24 |
| knapsack_handler_backtracking | | 163084480 | 163085075 | 4 |
| knapsack_handler_backtracking | | 559927626 | 559929100 | 1 |
| knapsack_handler_backtracking | | 1736294 | 1736248 | 410 |
| knapsack_handler_backtracking | | 18161163 | 18160888 | 40 |
| knapsack_handler_backtracking | | 89581115 | 89579544 | 9 |
| knapsack_handler_backtracking | | 3263062506 | 3263021100 | 1 |
| knapsack_handler_backtracking | | 3040190737 | 3040169900 | 1 |
| knapsack_handler_backtracking | | 211758763 | 211757767 | 3 |
| knapsack_handler_backtracking | | 624621926 | 624091200 | 1 |
| knapsack_handler_backtracking | | 1359326241 | 1359322300 | 1 |
| knapsack_handler_backtracking | | 72809409 | 72808410 | 10 |
| knapsack_handler_backtracking | | 6523479592 | 6523417800 | 1 |
| knapsack_handler_backtracking | | 342279551 | 342120750 | 2 |
| knapsack_handler_backtracking | | 1379948451 | 1379898500 | 1 |
| knapsack_handler_backtracking | | 125158735 | 125157260 | 5 |
| knapsack_handler_backtracking | | 5962969626 | 5962870200 | 1 |
| knapsack_handler_backtracking | | 34420000000 | 34420000000 | 1 |
| knapsack_handler_backtracking | | 123940051 | 123939083 | 6 |
| knapsack_handler_backtracking | | 31164000000 | 31163000000 | 1 |
| knapsack_handler_backtracking | | 47379010 | 47379250 | 14 |
| knapsack_handler_backtracking | | 2652311691 | 2652282900 | 1 |

Table 2. Benchmark Results for knapsack_handler_backtracking

| Benchmark Greedy | Time (ns) | CPU (ns) | Iterations |
|-------------------------|------------------|-----------------|-------------------|
| knapsack_handler_greedy | 7195 | 7193 | 93873 |
| knapsack_handler_greedy | 8058 | 8058 | 86586 |
| knapsack_handler_greedy | 9510 | 9510 | 77823 |
| knapsack_handler_greedy | 9905 | 9905 | 64998 |
| knapsack_handler_greedy | 9600 | 9600 | 72680 |
| knapsack_handler_greedy | 7909 | 7909 | 89973 |
| knapsack_handler_greedy | 13484 | 13484 | 52825 |
| knapsack_handler_greedy | 15715 | 15715 | 48553 |
| knapsack_handler_greedy | 16990 | 16990 | 45755 |
| knapsack_handler_greedy | 16546 | 16545 | 42428 |
| knapsack_handler_greedy | 18386 | 18366 | 34185 |
| knapsack_handler_greedy | 17087 | 17087 | 39253 |
| knapsack_handler_greedy | 18532 | 18528 | 38958 |
| knapsack_handler_greedy | 18495 | 18495 | 38447 |
| knapsack_handler_greedy | 20672 | 20671 | 32829 |
| knapsack_handler_greedy | 21150 | 21148 | 33215 |
| knapsack_handler_greedy | 21175 | 21175 | 33052 |
| knapsack_handler_greedy | 21954 | 21954 | 28701 |
| knapsack_handler_greedy | 23307 | 23307 | 34267 |
| knapsack_handler_greedy | 19612 | 19612 | 29183 |
| knapsack_handler_greedy | 21941 | 21938 | 36639 |
| knapsack_handler_greedy | 26130 | 26130 | 26768 |
| knapsack_handler_greedy | 28771 | 28741 | 24319 |
| knapsack_handler_greedy | 30160 | 30160 | 21867 |
| knapsack_handler_greedy | 30992 | 30991 | 22807 |
| knapsack_handler_greedy | 31482 | 31482 | 21237 |
| knapsack_handler_greedy | 34720 | 34719 | 21059 |
| knapsack_handler_greedy | 34224 | 34224 | 20451 |
| knapsack_handler_greedy | 7058 | 7058 | 90852 |
| knapsack_handler_greedy | 10121 | 10121 | 71612 |

Table 3. Benchmark Results for knapsack_handler_greedy

6 Conclusions

The following images are graphs representing the performance of each algorithm:




Fig. 1. Time taken for the Greedy benchmarks




Fig. 2. Time taken for the Backtracking benchmarks

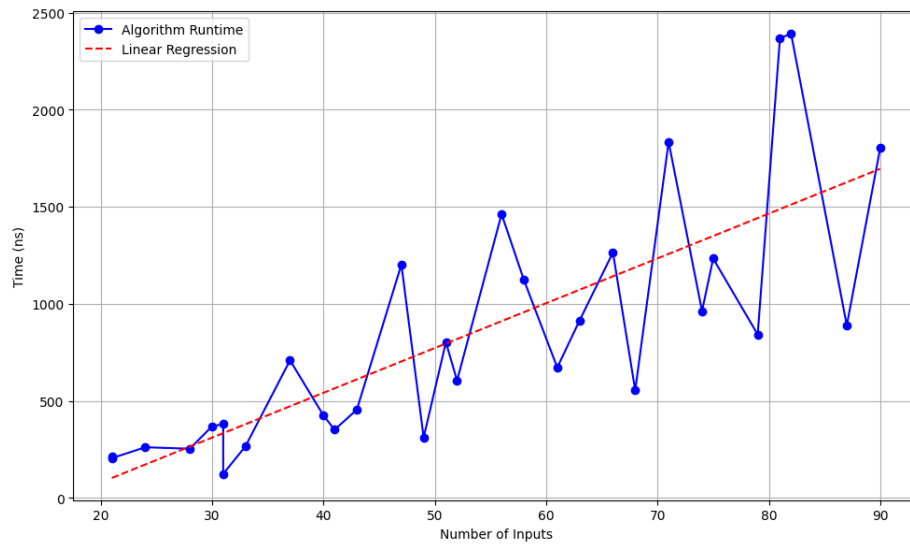


Fig. 3. Time taken for the Tabulation benchmarks

The graphs show the distribution of the performances of the algorithms based on the number of inputs and the execution time. Every algorithm shows an ascending trend, meaning that, generally, they take longer to complete the higher the number of sets of inputs they are given.

Some tests appear to be faster than others that have less sets of inputs. This is predominantly happening because the benchmarks were run in a multithreaded environment, where context switches were happening.

It is easy to say that the tabulation approach is by far the most reliable, albeit not the fastest. The speed of the greedy approach can be deceiving, suggesting it would be the best algorithm to solve the problem, but we have to take into consideration the fact that it does not provide the highest possible profit. Its solution is *almost* correct. With all of that in mind, if we are not looking for the perfect combination of items, but for a quick and decent result, the greedy would be the best algorithm. As seen in the second graph, the backtracking approach is by far the slowest of them all, since it explores every possible combination of items (time complexity of $\mathcal{O}(2^n)$). Even though it is very slow, it is still better in terms of the correctness of the result than the greedy algorithm. The only point it actually excels in is its simplicity of the implementation, as it's a good starting point to understand the problem.

With all of that being said, if we care about achieving a perfect result in a timely manner, we can conclude that the dynamic programming approach is the best one, and if we want a solution as fast as possible, although not perfect, the greedy algorithm is the way to go.

References

1. *Knapsack problem*, https://en.wikipedia.org/wiki/Knapsack_problem Accesat la [04/01/2025].
2. Kayode Badiru: *KNAPSACK PROBLEMS; METHODS, MODELS AND APPLICATIONS*, <https://shareok.org/server/api/core/bitstreams/7bb97e22-4c5d-44b3-b756-74aba3fc1fd7/content> Accesat la [05/01/2025]
3. LNCS Homepage, <http://www.springer.com/lncs>, Accesat la [04/01/2025]
4. A Phase Angle-Modulated Bat Algorithm with Application to Antenna Topology Optimization https://www.researchgate.net/publication/349878676_A_Phase_Angle-Modulated_Bat_Algorithm_with_Application_to_Antenna_Topology_Optimization, Accesat la [04/01/2025]